

EECS 391

Intro to AI

Evaluation Functions, Stochastic Games,
Other Approaches

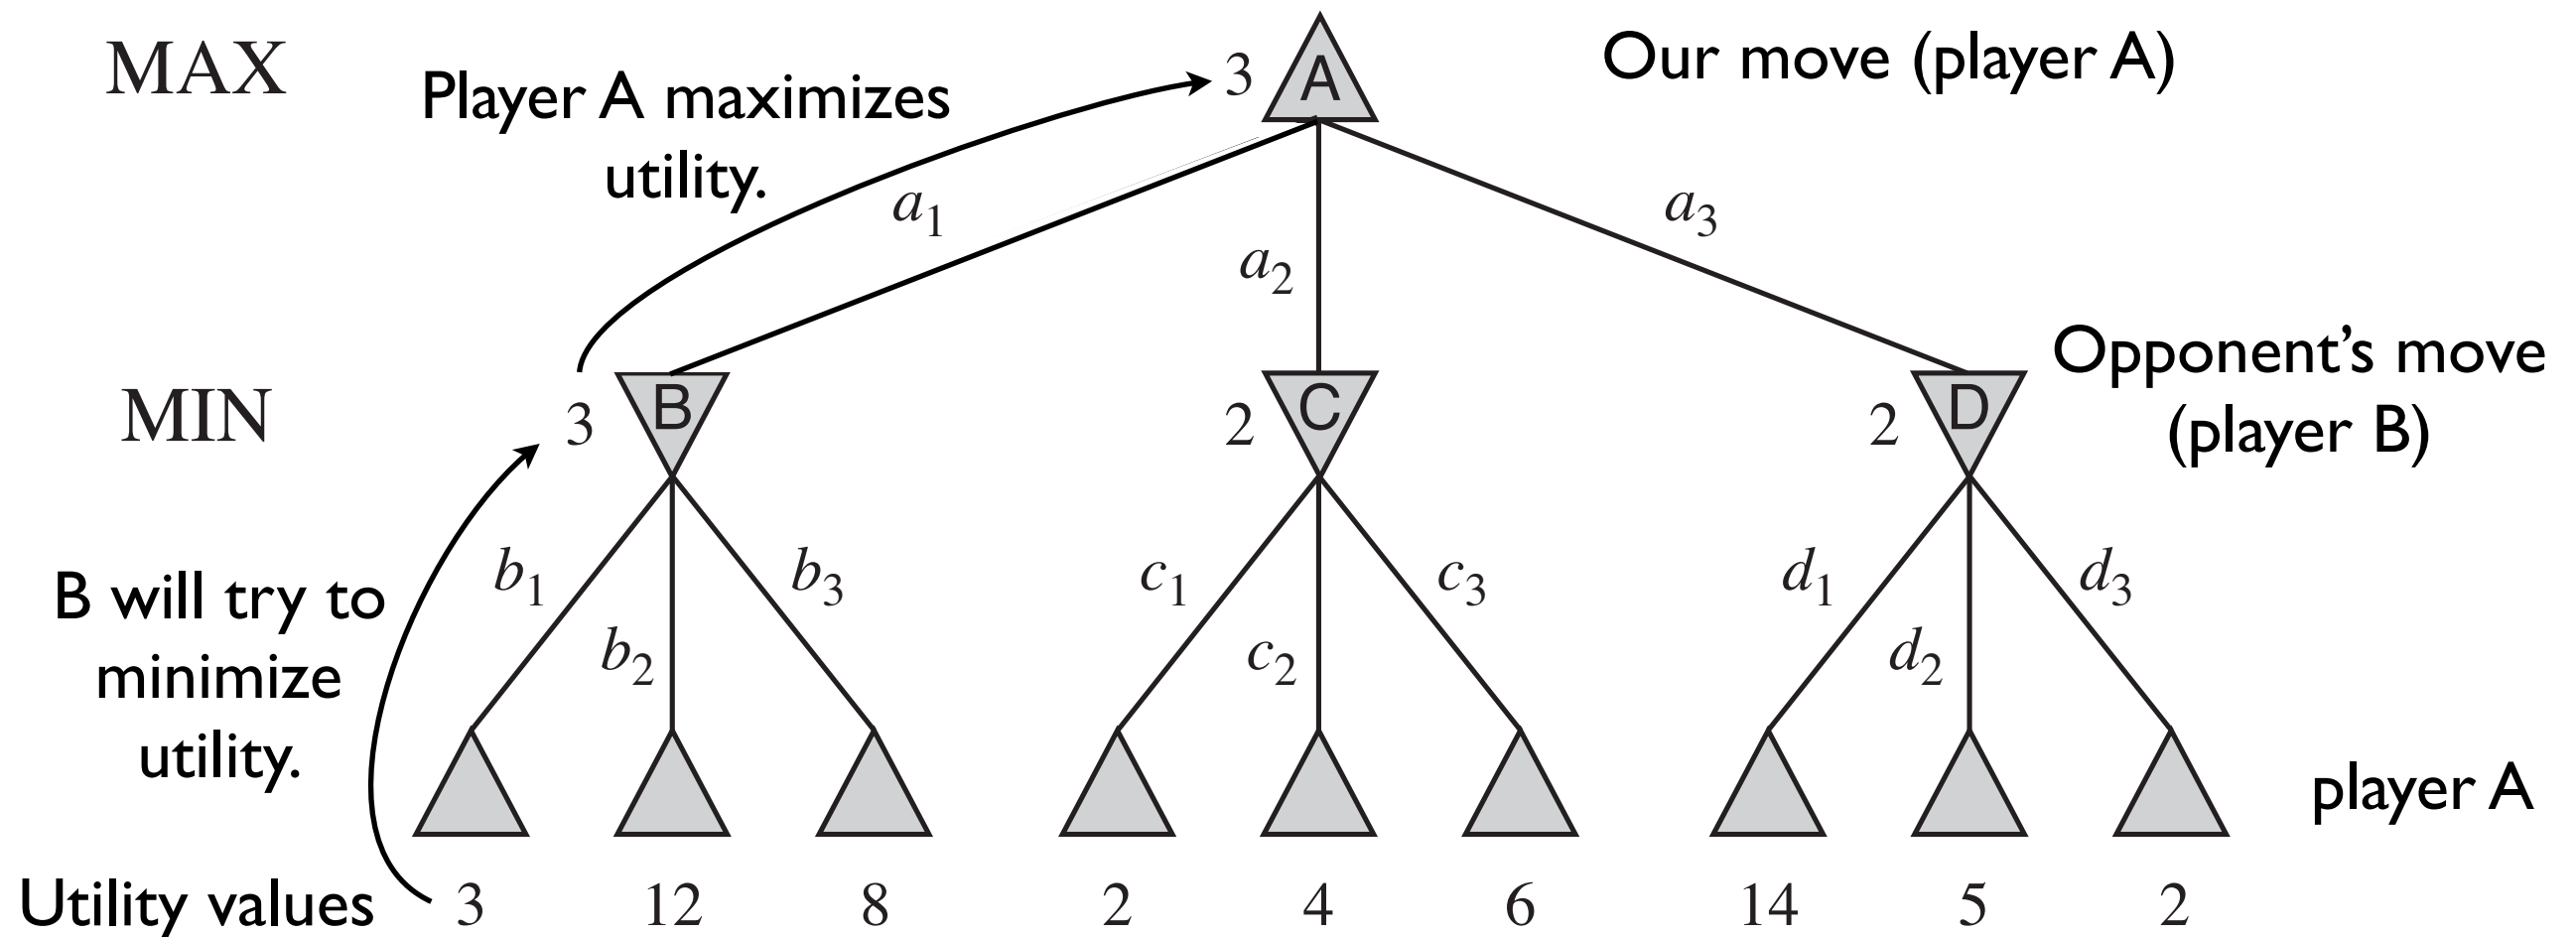
L7 Thu Sep 21, 2017

Imagine how great universities could be without all those human teachers

- Quartz series on “The Vanishing University”, a four-part series exploring the tech-driven future of higher education in America.
- *Jill Watson is the best damn teaching assistant you could ever want.*
- *“I got comments like ‘Mind blown’ and ‘I want to nominate Jill for outstanding TA award,’” says Ashok Goel, a computer-science professor at Georgia Tech.*
- Ashok Goel is also Jill Watson’s creator.
- Jill Watson is a chat bot.

Minimax: perfect play for deterministic 2-player games

- Idea: Chose move to position (or state) with highest minimax value.
Best achievable payoff against optimal player
- $\text{minimax-val}(n) =$
 - $\text{Utility}(n)$ if terminal state
 - $\max s \in \text{result}(n)$ if n is max node
 - $\min s \in \text{result}(n)$ if n is min node



Minimax algorithm

arg max returns the best action

function MINIMAX-DECISION(*state*) **returns** *an action*

return $\arg \max_{a \in \text{ACTIONS}(s)} \text{MIN-VALUE}(\text{RESULT}(state, a))$

transition model: state resulting from move *a* in state *state*

function MAX-VALUE(*state*) **returns** *a utility value*

if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow -\infty$

for each *a* **in** ACTIONS(*state*) **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a)))$

return *v*

Available actions from state *s*

Utility function: numerical value of state (defined only for terminal states)

function MIN-VALUE(*state*) **returns** *a utility value*

if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)


$v \leftarrow \infty$

for each *a* **in** ACTIONS(*state*) **do**

$v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a)))$

return *v*

Is the game over ?




Minimax algorithm

function MINIMAX-DECISION($state$) **returns** *an action*
return $\arg \max_{a \in \text{ACTIONS}(s)} \text{MIN-VALUE}(\text{RESULT}(\mathbf{A}, \mathbf{a}_1))$

function MAX-VALUE($state$) **returns** *a utility value*
if TERMINAL-TEST($state$) **then return** UTILITY($state$)
 $v \leftarrow -\infty$
for each a **in** ACTIONS($state$) **do**
 $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a)))$
return v

function MIN-VALUE($state$) **returns** *a utility value*
if TERMINAL-TEST($state$) **then return** UTILITY($state$)
 $v \leftarrow \infty$
for each a **in** ACTIONS($state$) **do**
 $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a)))$
return v




Minimax algorithm

function MINIMAX-DECISION(*state*) **returns** *an action*
return $\arg \max_{a \in \text{ACTIONS}(s)} \text{MIN-VALUE}(\text{RESULT}(\mathbf{A}, \mathbf{a}_1))$

function MAX-VALUE(*state*) **returns** *a utility value*
if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)
 $v \leftarrow -\infty$
for each *a* **in** ACTIONS(*state*) **do**
 $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a)))$
return *v*

function MIN-VALUE(\mathbf{B}) **returns** *a utility value*
if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)
 $v \leftarrow \infty$
for each *a* **in** ACTIONS(*state*) **do**
 $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a)))$
return *v*




Minimax algorithm

function MINIMAX-DECISION(*state*) **returns** *an action*
return $\arg \max_{a \in \text{ACTIONS}(s)} \text{MIN-VALUE}(\text{RESULT}(\mathbf{A}, \mathbf{a}_1))$


function MAX-VALUE(*state*) **returns** *a utility value*
if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)
 $v \leftarrow -\infty$
for each *a* **in** ACTIONS(*state*) **do**
 $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a)))$
return *v*

function MIN-VALUE(\mathbf{B}) **returns** *a utility value*
if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)
 $v \leftarrow \infty$
for each *a* **in** ACTIONS(*state*) **do**
 $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(\mathbf{B}, \mathbf{b}_1)))$
return *v*



Alpha-beta pruning


- Idea: prune branches that can't influence final decision
- alpha-beta pruning keeps track of the range of possible values:
 α = max's best choice (highest val.) so far at any point along path
 β = min's best choice (lowest val.) so far at any point along path
- alpha-beta doesn't search ("*prunes*"), branches that don't affect current estimates



Minimax: perfect play for deterministic 2-player games


- $\text{minimax}(n) =$
 - $\text{Utility}(n)$ if terminal state
 - $\max s \in \text{result}(n)$ if n is max node
 - $\min s \in \text{result}(n)$ if n is min node
- $\text{minimax}(A) = \max(\min(3, 12, 8), \min(2, 4, 6), \min(14, 5, 2))$

Utilities are computed recursively from terminal states.




Minimax: perfect play for deterministic 2-player games

- $\text{minimax}(n) =$
 - $\text{Utility}(n)$ if terminal state
 - $\max s \in \text{result}(n)$ if n is max node
 - $\min s \in \text{result}(n)$ if n is min node
- $\text{minimax}(A) = \max(\min(3, 12, 8), \min(2, x, y), \min(14, 5, 2))$




Minimax: perfect play for deterministic 2-player games

- $\text{minimax}(n) =$
 - $\text{Utility}(n)$ if terminal state
 - $\max s \in \text{result}(n)$ if n is max node
 - $\min s \in \text{result}(n)$ if n is min node
- $\text{minimax}(A) = \max(\min(3, 12, 8), \min(2, -5, -9), \min(14, 5, 2))$
 $= -9$
but $\max(3, -9, 2) = 3 \Rightarrow a1$



Minimax: perfect play for deterministic 2-player games

- $\text{minimax}(n) =$
 - $\text{Utility}(n)$ if terminal state
 - $\max s \in \text{result}(n)$ if n is max node
 - $\min s \in \text{result}(n)$ if n is min node
- $\text{minimax}(A) = \max(\min(3, 12, 8), \min(2, -5, -9), \min(14, 5, 2))$
 $= -9$
but $\max(3, -9, 2) = 3 \Rightarrow a_1$



Alpha-Beta Search Algorithm

function ALPHA-BETA-SEARCH(*state*) **returns** an action
 $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$
 return the *action* in ACTIONS(*state*) with value v

Alpha-beta pruning is a modification of the standard mini-max algorithm.

function MAX-VALUE(*state*, α , β) **returns** a utility value
 if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)
 $v \leftarrow -\infty$
 for each a **in** ACTIONS(*state*) **do**
 $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$
 if $v \geq \beta$ **then return** v
 $\alpha \leftarrow \text{MAX}(\alpha, v)$
 return v

function MIN-VALUE(*state*, α , β) **returns** a utility value
 if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)
 $v \leftarrow +\infty$
 for each a **in** ACTIONS(*state*) **do**
 $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$
 if $v \leq \alpha$ **then return** v
 $\beta \leftarrow \text{MIN}(\beta, v)$
 return v

It cuts off the search, when the result can't change.

Alpha-Beta Search Algorithm

function ALPHA-BETA-SEARCH($state$) **returns** an action
 $v \leftarrow \text{MAX-VALUE}(state, -\infty, +\infty)$
 return the *action* in ACTIONS($state$) with value v

Alpha-beta pruning is a modification of the standard mini-max algorithm.


function MAX-VALUE($state, \alpha, \beta$) **returns** a utility value
 if TERMINAL-TEST($state$) **then return** UTILITY($state$)
 $v \leftarrow -\infty$
 for each a **in** ACTIONS($state$) **do**
 $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$
 if $v \geq \beta$ **then return** v
 $\alpha \leftarrow \text{MAX}(\alpha, v)$
 return v

function MIN-VALUE($state, \alpha, \beta$) **returns** a utility value
 if TERMINAL-TEST($state$) **then return** UTILITY($state$)
 $v \leftarrow +\infty$
 for each a **in** ACTIONS($state$) **do**
 $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$
 if $v \leq \alpha$ **then return** v
 $\beta \leftarrow \text{MIN}(\beta, v)$
 return v

Otherwise, it updates the best value found so far.

More properties of alpha-beta pruning

- applies to tree of any depth
- can prune entire sub-trees, not just leaves
- Idea:
 - If node n has greater value than node m , then
 - *node n would never be reached*
 - because our Opponent (the min player) would have to choose node m to be optimal
- That means the entire branch can be pruned once we find such a node.




Alpha-beta pruning

- Idea: prune branches that can't influence final decision
- alpha-beta pruning keeps track of the range of possible values:

α = max's best choice (highest val.) so far at any point along path

β = min's best choice (lowest val.) so far at any point along path

- alpha-beta doesn't search: it “prunes” branches that don't affect current estimates



```
function ALPHA-BETA-SEARCH(state) returns an action  
   $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$   
  return the action in ACTIONS(state) with value v
```

```
function MAX-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow -\infty$   
  for each a in ACTIONS(state) do  
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \geq \beta$  then return v  
     $\alpha \leftarrow \text{MAX}(\alpha, v)$   
  return v
```

```
function MIN-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow +\infty$   
  for each a in ACTIONS(state) do  
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \leq \alpha$  then return v  
     $\beta \leftarrow \text{MIN}(\beta, v)$   
  return v
```



Alpha-beta pruning

- Idea: prune branches that can't influence final decision
- alpha-beta pruning keeps track of the range of possible values:

α = max's best choice (highest val.) so far at any point along path

β = min's best choice (lowest val.) so far at any point along path

- alpha-beta doesn't search: it “prunes” branches that don't affect current estimates



```
function ALPHA-BETA-SEARCH(state) returns an action  
   $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$   
  return the action in  $\text{ACTIONS}(\text{state})$  with value  $v$ 
```

```
function MAX-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
if  $\text{TERMINAL-TEST}(\text{state})$  then return  $\text{UTILITY}(\text{state})$   
   $v \leftarrow -\infty$   
  for each  $a$  in  $\text{ACTIONS}(\text{state})$  do  
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \geq \beta$  then return  $v$   
     $\alpha \leftarrow \text{MAX}(\alpha, v)$   
  return  $v$ 
```

```
function MIN-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
if  $\text{TERMINAL-TEST}(\text{state})$  then return  $\text{UTILITY}(\text{state})$   
   $v \leftarrow +\infty$   
  for each  $a$  in  $\text{ACTIONS}(\text{state})$  do  
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \leq \alpha$  then return  $v$   
     $\beta \leftarrow \text{MIN}(\beta, v)$   
  return  $v$ 
```


Alpha-beta pruning

- Idea: prune branches that can't influence final decision
- alpha-beta pruning keeps track of the range of possible values:

α = max's best choice (highest val.) so far at any point along path

β = min's best choice (lowest val.) so far at any point along path

- alpha-beta doesn't search: it “prunes” branches that don't affect current estimates



```
function ALPHA-BETA-SEARCH(state) returns an action  
   $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$   
  return the action in ACTIONS(state) with value v
```

```
function MAX-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow -\infty$   
  for each a in ACTIONS(state) do  
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \geq \beta$  then return v  
     $\alpha \leftarrow \text{MAX}(\alpha, v)$   
  return v
```

```
function MIN-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow +\infty$   
  for each a in ACTIONS(state) do  
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \leq \alpha$  then return v  
     $\beta \leftarrow \text{MIN}(\beta, v)$   
  return v
```


Alpha-beta pruning

- Idea: prune branches that can't influence final decision
- alpha-beta pruning keeps track of the range of possible values:

α = max's best choice (highest val.) so far at any point along path

β = min's best choice (lowest val.) so far at any point along path

- alpha-beta doesn't search: it “prunes” branches that don't affect current estimates



```
function ALPHA-BETA-SEARCH(state) returns an action  
   $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$   
  return the action in ACTIONS(state) with value v
```

```
function MAX-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow -\infty$   
  for each a in ACTIONS(state) do  
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \geq \beta$  then return v  
     $\alpha \leftarrow \text{MAX}(\alpha, v)$   
  return v
```

```
function MIN-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow +\infty$   
  for each a in ACTIONS(state) do  
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \leq \alpha$  then return v  
     $\beta \leftarrow \text{MIN}(\beta, v)$   
  return v
```


Alpha-beta pruning

- Idea: prune branches that can't influence final decision
- alpha-beta pruning keeps track of the range of possible values:

α = max's best choice (highest val.) so far at any point along path

β = min's best choice (lowest val.) so far at any point along path

- alpha-beta doesn't search: it “prunes” branches that don't affect current estimates



```
function ALPHA-BETA-SEARCH(state) returns an action  
   $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$   
  return the action in ACTIONS(state) with value  $v$ 
```

```
function MAX-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow -\infty$   
  for each  $a$  in ACTIONS(state) do  
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \geq \beta$  then return  $v$   
     $\alpha \leftarrow \text{MAX}(\alpha, v)$   
  return  $v$ 
```

```
function MIN-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow +\infty$   
  for each  $a$  in ACTIONS(state) do  
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \leq \alpha$  then return  $v$   
     $\beta \leftarrow \text{MIN}(\beta, v)$   
  return  $v$ 
```


Alpha-beta pruning

- Idea: prune branches that can't influence final decision
- alpha-beta pruning keeps track of the range of possible values:

α = max's best choice (highest val.) so far at any point along path

β = min's best choice (lowest val.) so far at any point along path

- alpha-beta doesn't search: it “prunes” branches that don't affect current estimates



```
function ALPHA-BETA-SEARCH(state) returns an action  
   $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$   
  return the action in ACTIONS(state) with value v
```

```
function MAX-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow -\infty$   
  for each a in ACTIONS(state) do  
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \geq \beta$  then return v  
     $\alpha \leftarrow \text{MAX}(\alpha, v)$   
  return v
```

```
function MIN-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow +\infty$   
  for each a in ACTIONS(state) do  
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \leq \alpha$  then return v  
     $\beta \leftarrow \text{MIN}(\beta, v)$   
  return v
```


Alpha-beta pruning

- Idea: prune branches that can't influence final decision
- alpha-beta pruning keeps track of the range of possible values:

α = max's best choice (highest val.) so far at any point along path

β = min's best choice (lowest val.) so far at any point along path

- alpha-beta doesn't search: it “prunes” branches that don't affect current estimates



```
function ALPHA-BETA-SEARCH(state) returns an action  
   $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$   
  return the action in ACTIONS(state) with value v
```

```
function MAX-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow -\infty$   
  for each a in ACTIONS(state) do  
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \geq \beta$  then return v  
     $\alpha \leftarrow \text{MAX}(\alpha, v)$   
  return v
```

```
function MIN-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow +\infty$   
  for each a in ACTIONS(state) do  
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \leq \alpha$  then return v  
     $\beta \leftarrow \text{MIN}(\beta, v)$   
  return v
```


Alpha-beta pruning

- Idea: prune branches that can't influence final decision
- alpha-beta pruning keeps track of the range of possible values:

α = max's best choice (highest val.) so far at any point along path

β = min's best choice (lowest val.) so far at any point along path

- alpha-beta doesn't search: it “prunes” branches that don't affect current estimates



```
function ALPHA-BETA-SEARCH(state) returns an action  
   $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$   
  return the action in ACTIONS(state) with value v
```

```
function MAX-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow -\infty$   
  for each a in ACTIONS(state) do  
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \geq \beta$  then return v  
     $\alpha \leftarrow \text{MAX}(\alpha, v)$   
  return v
```

```
function MIN-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow +\infty$   
  for each a in ACTIONS(state) do  
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \leq \alpha$  then return v  
     $\beta \leftarrow \text{MIN}(\beta, v)$   
  return v
```



Alpha-beta pruning

- Idea: prune branches that can't influence final decision
- alpha-beta pruning keeps track of the range of possible values:

α = max's best choice (highest val.) so far at any point along path

β = min's best choice (lowest val.) so far at any point along path

- alpha-beta doesn't search: it “prunes” branches that don't affect current estimates



```
function ALPHA-BETA-SEARCH(state) returns an action
   $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$ 
  return the action in ACTIONS(state) with value  $v$ 
```

```
function MAX-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value
  if TERMINAL-TEST(state) then return UTILITY(state)
   $v \leftarrow -\infty$ 
  for each  $a$  in ACTIONS(state) do
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$ 
    if  $v \geq \beta$  then return  $v$ 
     $\alpha \leftarrow \text{MAX}(\alpha, v)$ 
  return  $v$ 
```

```
function MIN-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value
  if TERMINAL-TEST(state) then return UTILITY(state)
   $v \leftarrow +\infty$ 
  for each  $a$  in ACTIONS(state) do
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$ 
    if  $v \leq \alpha$  then return  $v$ 
     $\beta \leftarrow \text{MIN}(\beta, v)$ 
  return  $v$ 
```



Alpha-beta pruning

- Idea: prune branches that can't influence final decision
- alpha-beta pruning keeps track of the range of possible values:

α = max's best choice (highest val.) so far at any point along path

β = min's best choice (lowest val.) so far at any point along path

- alpha-beta doesn't search: it “prunes” branches that don't affect current estimates



```
function ALPHA-BETA-SEARCH(state) returns an action  
   $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$   
  return the action in  $\text{ACTIONS}(\text{state})$  with value  $v$ 
```

```
function MAX-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if  $\text{TERMINAL-TEST}(\text{state})$  then return  $\text{UTILITY}(\text{state})$   
   $v \leftarrow -\infty$   
  for each  $a$  in  $\text{ACTIONS}(\text{state})$  do  
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \geq \beta$  then return  $v$   
     $\alpha \leftarrow \text{MAX}(\alpha, v)$   
  return  $v$ 
```

```
function MIN-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if  $\text{TERMINAL-TEST}(\text{state})$  then return  $\text{UTILITY}(\text{state})$   
   $v \leftarrow +\infty$   
  for each  $a$  in  $\text{ACTIONS}(\text{state})$  do  
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \leq \alpha$  then return  $v$   
     $\beta \leftarrow \text{MIN}(\beta, v)$   
  return  $v$ 
```


Alpha-beta pruning

- Idea: prune branches that can't influence final decision
- alpha-beta pruning keeps track of the range of possible values:

α = max's best choice (highest val.) so far at any point along path

β = min's best choice (lowest val.) so far at any point along path

- alpha-beta doesn't search: it “prunes” branches that don't affect current estimates



```
function ALPHA-BETA-SEARCH(state) returns an action  
   $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$   
  return the action in ACTIONS(state) with value v
```

```
function MAX-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow -\infty$   
  for each a in ACTIONS(state) do  
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \geq \beta$  then return v  
     $\alpha \leftarrow \text{MAX}(\alpha, v)$   
  return v
```

```
function MIN-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow +\infty$   
  for each a in ACTIONS(state) do  
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \leq \alpha$  then return v  
     $\beta \leftarrow \text{MIN}(\beta, v)$   
  return v
```


Alpha-beta pruning

- Idea: prune branches that can't influence final decision
- alpha-beta pruning keeps track of the range of possible values:

α = max's best choice (highest val.) so far at any point along path

β = min's best choice (lowest val.) so far at any point along path

- alpha-beta doesn't search: it “prunes” branches that don't affect current estimates



```
function ALPHA-BETA-SEARCH(state) returns an action  
   $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$   
  return the action in ACTIONS(state) with value v
```

```
function MAX-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow -\infty$   
  for each a in ACTIONS(state) do  
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \geq \beta$  then return v  
     $\alpha \leftarrow \text{MAX}(\alpha, v)$   
  return v
```

```
function MIN-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow +\infty$   
  for each a in ACTIONS(state) do  
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \leq \alpha$  then return v  
     $\beta \leftarrow \text{MIN}(\beta, v)$   
  return v
```


Alpha-beta pruning

- Idea: prune branches that can't influence final decision
- alpha-beta pruning keeps track of the range of possible values:

α = max's best choice (highest val.) so far at any point along path

β = min's best choice (lowest val.) so far at any point along path

- alpha-beta doesn't search: it “prunes” branches that don't affect current estimates



```
function ALPHA-BETA-SEARCH(state) returns an action  
   $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$   
  return the action in  $\text{ACTIONS}(\text{state})$  with value  $v$ 
```

```
function MAX-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if  $\text{TERMINAL-TEST}(\text{state})$  then return  $\text{UTILITY}(\text{state})$   
   $v \leftarrow -\infty$   
  for each  $a$  in  $\text{ACTIONS}(\text{state})$  do  
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \geq \beta$  then return  $v$   
     $\alpha \leftarrow \text{MAX}(\alpha, v)$   
  return  $v$ 
```

```
function MIN-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if  $\text{TERMINAL-TEST}(\text{state})$  then return  $\text{UTILITY}(\text{state})$   
   $v \leftarrow +\infty$   
  for each  $a$  in  $\text{ACTIONS}(\text{state})$  do  
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \leq \alpha$  then return  $v$   
     $\beta \leftarrow \text{MIN}(\beta, v)$   
  return  $v$ 
```


Alpha-beta pruning

- Idea: prune branches that can't influence final decision
- alpha-beta pruning keeps track of the range of possible values:

α = max's best choice (highest val.) so far at any point along path

β = min's best choice (lowest val.) so far at any point along path

- alpha-beta doesn't search: it “prunes” branches that don't affect current estimates



```
function ALPHA-BETA-SEARCH(state) returns an action  
   $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$   
  return the action in  $\text{ACTIONS}(\text{state})$  with value  $v$ 
```

```
function MAX-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if  $\text{TERMINAL-TEST}(\text{state})$  then return  $\text{UTILITY}(\text{state})$   
   $v \leftarrow -\infty$   
  for each  $a$  in  $\text{ACTIONS}(\text{state})$  do  
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \geq \beta$  then return  $v$   
     $\alpha \leftarrow \text{MAX}(\alpha, v)$   
  return  $v$ 
```

```
function MIN-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if  $\text{TERMINAL-TEST}(\text{state})$  then return  $\text{UTILITY}(\text{state})$   
   $v \leftarrow +\infty$   
  for each  $a$  in  $\text{ACTIONS}(\text{state})$  do  
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \leq \alpha$  then return  $v$   
     $\beta \leftarrow \text{MIN}(\beta, v)$   
  return  $v$ 
```


Alpha-beta pruning

- Idea: prune branches that can't influence final decision
- alpha-beta pruning keeps track of the range of possible values:

α = max's best choice (highest val.) so far at any point along path

β = min's best choice (lowest val.) so far at any point along path

- alpha-beta doesn't search: it “prunes” branches that don't affect current estimates



```
function ALPHA-BETA-SEARCH(state) returns an action  
   $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$   
  return the action in ACTIONS(state) with value  $v$ 
```

```
function MAX-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow -\infty$   
  for each  $a$  in ACTIONS(state) do  
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \geq \beta$  then return  $v$   
     $\alpha \leftarrow \text{MAX}(\alpha, v)$   
  return  $v$ 
```

```
function MIN-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow +\infty$   
  for each  $a$  in ACTIONS(state) do  
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \leq \alpha$  then return  $v$   
     $\beta \leftarrow \text{MIN}(\beta, v)$   
  return  $v$ 
```


Alpha-beta pruning

- Idea: prune branches that can't influence final decision
- alpha-beta pruning keeps track of the range of possible values:

α = max's best choice (highest val.) so far at any point along path

β = min's best choice (lowest val.) so far at any point along path

- alpha-beta doesn't search: it “prunes” branches that don't affect current estimates



```
function ALPHA-BETA-SEARCH(state) returns an action  
   $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$   
  return the action in  $\text{ACTIONS}(\text{state})$  with value  $v$ 
```

```
function MAX-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if  $\text{TERMINAL-TEST}(\text{state})$  then return  $\text{UTILITY}(\text{state})$   
   $v \leftarrow -\infty$   
  for each  $a$  in  $\text{ACTIONS}(\text{state})$  do  
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \geq \beta$  then return  $v$   
     $\alpha \leftarrow \text{MAX}(\alpha, v)$   
  return  $v$ 
```

```
function MIN-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if  $\text{TERMINAL-TEST}(\text{state})$  then return  $\text{UTILITY}(\text{state})$   
   $v \leftarrow +\infty$   
  for each  $a$  in  $\text{ACTIONS}(\text{state})$  do  
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \leq \alpha$  then return  $v$   
     $\beta \leftarrow \text{MIN}(\beta, v)$   
  return  $v$ 
```



Alpha-beta pruning

- Idea: prune branches that can't influence final decision
- alpha-beta pruning keeps track of the range of possible values:

α = max's best choice (highest val.) so far at any point along path

β = min's best choice (lowest val.) so far at any point along path

- alpha-beta doesn't search: it “prunes” branches that don't affect current estimates



```
function ALPHA-BETA-SEARCH(state) returns an action  
   $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$   
  return the action in  $\text{ACTIONS}(\text{state})$  with value  $v$ 
```

```
function MAX-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if  $\text{TERMINAL-TEST}(\text{state})$  then return  $\text{UTILITY}(\text{state})$   
   $v \leftarrow -\infty$   
  for each  $a$  in  $\text{ACTIONS}(\text{state})$  do  
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \geq \beta$  then return  $v$   
     $\alpha \leftarrow \text{MAX}(\alpha, v)$   
  return  $v$ 
```

```
function MIN-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if  $\text{TERMINAL-TEST}(\text{state})$  then return  $\text{UTILITY}(\text{state})$   
   $v \leftarrow +\infty$   
  for each  $a$  in  $\text{ACTIONS}(\text{state})$  do  
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \leq \alpha$  then return  $v$   
     $\beta \leftarrow \text{MIN}(\beta, v)$   
  return  $v$ 
```



Alpha-beta pruning

- Idea: prune branches that can't influence final decision
- alpha-beta pruning keeps track of the range of possible values:

α = max's best choice (highest val.) so far at any point along path

β = min's best choice (lowest val.) so far at any point along path

- alpha-beta doesn't search: it “prunes” branches that don't affect current estimates



```
function ALPHA-BETA-SEARCH(state) returns an action
   $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$ 
  return the action in  $\text{ACTIONS}(\text{state})$  with value  $v$ 
```

```
function MAX-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value
  if  $\text{TERMINAL-TEST}(\text{state})$  then return  $\text{UTILITY}(\text{state})$ 
   $v \leftarrow -\infty$ 
  for each  $a$  in  $\text{ACTIONS}(\text{state})$  do
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$ 
    if  $v \geq \beta$  then return  $v$ 
     $\alpha \leftarrow \text{MAX}(\alpha, v)$ 
  return  $v$ 
```

```
function MIN-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value
  if  $\text{TERMINAL-TEST}(\text{state})$  then return  $\text{UTILITY}(\text{state})$ 
   $v \leftarrow +\infty$ 
  for each  $a$  in  $\text{ACTIONS}(\text{state})$  do
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$ 
    if  $v \leq \alpha$  then return  $v$ 
     $\beta \leftarrow \text{MIN}(\beta, v)$ 
  return  $v$ 
```


Alpha-beta pruning

- Idea: prune branches that can't influence final decision
- alpha-beta pruning keeps track of the range of possible values:

α = max's best choice (highest val.) so far at any point along path

β = min's best choice (lowest val.) so far at any point along path

- alpha-beta doesn't search: it “prunes” branches that don't affect current estimates



```
function ALPHA-BETA-SEARCH(state) returns an action
   $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$ 
  return the action in ACTIONS(state) with value v
```

```
function MAX-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value
  if TERMINAL-TEST(state) then return UTILITY(state)
   $v \leftarrow -\infty$ 
  for each a in ACTIONS(state) do
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$ 
    if  $v \geq \beta$  then return v
     $\alpha \leftarrow \text{MAX}(\alpha, v)$ 
  return v
```

```
function MIN-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value
  if TERMINAL-TEST(state) then return UTILITY(state)
   $v \leftarrow +\infty$ 
  for each a in ACTIONS(state) do
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$ 
    if  $v \leq \alpha$  then return v
     $\beta \leftarrow \text{MIN}(\beta, v)$ 
  return v
```


Alpha-beta pruning

- Idea: prune branches that can't influence final decision
- alpha-beta pruning keeps track of the range of possible values:

α = max's best choice (highest val.) so far at any point along path

β = min's best choice (lowest val.) so far at any point along path

- alpha-beta doesn't search: it “prunes” branches that don't affect current estimates



```
function ALPHA-BETA-SEARCH(state) returns an action  
   $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$   
  return the action in  $\text{ACTIONS}(\text{state})$  with value  $v$ 
```

```
function MAX-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if  $\text{TERMINAL-TEST}(\text{state})$  then return  $\text{UTILITY}(\text{state})$   
   $v \leftarrow -\infty$   
  for each  $a$  in  $\text{ACTIONS}(\text{state})$  do  
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \geq \beta$  then return  $v$   
     $\alpha \leftarrow \text{MAX}(\alpha, v)$   
  return  $v$ 
```

```
function MIN-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if  $\text{TERMINAL-TEST}(\text{state})$  then return  $\text{UTILITY}(\text{state})$   
   $v \leftarrow +\infty$   
  for each  $a$  in  $\text{ACTIONS}(\text{state})$  do  
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \leq \alpha$  then return  $v$   
     $\beta \leftarrow \text{MIN}(\beta, v)$   
  return  $v$ 
```


Alpha-beta pruning

- Idea: prune branches that can't influence final decision
- alpha-beta pruning keeps track of the range of possible values:

α = max's best choice (highest val.) so far at any point along path

β = min's best choice (lowest val.) so far at any point along path

- alpha-beta doesn't search: it “prunes” branches that don't affect current estimates



```
function ALPHA-BETA-SEARCH(state) returns an action  
   $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$   
  return the action in ACTIONS(state) with value  $v$ 
```

```
function MAX-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow -\infty$   
  for each  $a$  in ACTIONS(state) do  
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \geq \beta$  then return  $v$   
     $\alpha \leftarrow \text{MAX}(\alpha, v)$   
  return  $v$ 
```

```
function MIN-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow +\infty$   
  for each  $a$  in ACTIONS(state) do  
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \leq \alpha$  then return  $v$   
     $\beta \leftarrow \text{MIN}(\beta, v)$   
  return  $v$ 
```


Alpha-beta pruning

- Idea: prune branches that can't influence final decision
- alpha-beta pruning keeps track of the range of possible values:

α = max's best choice (highest val.) so far at any point along path

β = min's best choice (lowest val.) so far at any point along path

- alpha-beta doesn't search: it “prunes” branches that don't affect current estimates



```
function ALPHA-BETA-SEARCH(state) returns an action  
   $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$   
  return the action in ACTIONS(state) with value  $v$ 
```

```
function MAX-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow -\infty$   
  for each  $a$  in ACTIONS(state) do  
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \geq \beta$  then return  $v$   
     $\alpha \leftarrow \text{MAX}(\alpha, v)$   
  return  $v$ 
```

```
function MIN-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow +\infty$   
  for each  $a$  in ACTIONS(state) do  
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \leq \alpha$  then return  $v$   
     $\beta \leftarrow \text{MIN}(\beta, v)$   
  return  $v$ 
```


Alpha-beta pruning

- Idea: prune branches that can't influence final decision
- alpha-beta pruning keeps track of the range of possible values:

α = max's best choice (highest val.) so far at any point along path

β = min's best choice (lowest val.) so far at any point along path

- alpha-beta doesn't search: it “prunes” branches that don't affect current estimates



```
function ALPHA-BETA-SEARCH(state) returns an action  
   $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$   
  return the action in ACTIONS(state) with value v
```

```
function MAX-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow -\infty$   
  for each a in ACTIONS(state) do  
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \geq \beta$  then return v  
     $\alpha \leftarrow \text{MAX}(\alpha, v)$   
  return v
```

```
function MIN-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow +\infty$   
  for each a in ACTIONS(state) do  
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \leq \alpha$  then return v  
     $\beta \leftarrow \text{MIN}(\beta, v)$   
  return v
```


Alpha-beta pruning

- Idea: prune branches that can't influence final decision
- alpha-beta pruning keeps track of the range of possible values:

α = max's best choice (highest val.) so far at any point along path

β = min's best choice (lowest val.) so far at any point along path

- alpha-beta doesn't search: it “prunes” branches that don't affect current estimates



```
function ALPHA-BETA-SEARCH(state) returns an action  
   $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$   
  return the action in  $\text{ACTIONS}(\text{state})$  with value  $v$ 
```

```
function MAX-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if  $\text{TERMINAL-TEST}(\text{state})$  then return  $\text{UTILITY}(\text{state})$   
   $v \leftarrow -\infty$   
  for each  $a$  in  $\text{ACTIONS}(\text{state})$  do  
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \geq \beta$  then return  $v$   
     $\alpha \leftarrow \text{MAX}(\alpha, v)$   
  return  $v$ 
```

```
function MIN-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if  $\text{TERMINAL-TEST}(\text{state})$  then return  $\text{UTILITY}(\text{state})$   
   $v \leftarrow +\infty$   
  for each  $a$  in  $\text{ACTIONS}(\text{state})$  do  
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \leq \alpha$  then return  $v$   
     $\beta \leftarrow \text{MIN}(\beta, v)$   
  return  $v$ 
```



Alpha-beta pruning

- Idea: prune branches that can't influence final decision
- alpha-beta pruning keeps track of the range of possible values:

α = max's best choice (highest val.) so far at any point along path

β = min's best choice (lowest val.) so far at any point along path

- alpha-beta doesn't search: it “prunes” branches that don't affect current estimates



```
function ALPHA-BETA-SEARCH(state) returns an action
   $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$ 
  return the action in ACTIONS(state) with value  $v$ 
```

```
function MAX-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value
  if TERMINAL-TEST(state) then return UTILITY(state)
   $v \leftarrow -\infty$ 
  for each  $a$  in ACTIONS(state) do
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$ 
    if  $v \geq \beta$  then return  $v$ 
     $\alpha \leftarrow \text{MAX}(\alpha, v)$ 
  return  $v$ 
```

```
function MIN-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value
  if TERMINAL-TEST(state) then return UTILITY(state)
   $v \leftarrow +\infty$ 
  for each  $a$  in ACTIONS(state) do
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$ 
    if  $v \leq \alpha$  then return  $v$ 
     $\beta \leftarrow \text{MIN}(\beta, v)$ 
  return  $v$ 
```


Alpha-beta pruning


- Idea: prune branches that can't influence final decision
- alpha-beta pruning keeps track of the range of possible values:

α = max's best choice (highest val.) so far at any point along path

β = min's best choice (lowest val.) so far at any point along path

- alpha-beta doesn't search: it “prunes” branches that don't affect current estimates

The *most* this branch can be is 2, and max already has a choice with value 3.




```
function ALPHA-BETA-SEARCH(state) returns an action  
   $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$   
  return the action in  $\text{ACTIONS}(\text{state})$  with value  $v$ 
```

```
function MAX-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if  $\text{TERMINAL-TEST}(\text{state})$  then return  $\text{UTILITY}(\text{state})$   
   $v \leftarrow -\infty$   
  for each  $a$  in  $\text{ACTIONS}(\text{state})$  do  
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \geq \beta$  then return  $v$   
     $\alpha \leftarrow \text{MAX}(\alpha, v)$   
  return  $v$ 
```


```
function MIN-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if  $\text{TERMINAL-TEST}(\text{state})$  then return  $\text{UTILITY}(\text{state})$   
   $v \leftarrow +\infty$   
  for each  $a$  in  $\text{ACTIONS}(\text{state})$  do  
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \leq \alpha$  then return  $v$   
     $\beta \leftarrow \text{MIN}(\beta, v)$   
  return  $v$ 
```

More properties of alpha-beta pruning


- can prune entire sub-trees, not just leaves
- Can result in a huge savings, e.g. $O(b^{m/2})$ vs $O(b^m)$
- search trees twice as deep
- But chess is still $O(35^{50})$
- *Problem*: game trees are far too deep.
- *Idea*: estimate *expected* utility based on game position evaluation function.





evaluation function example



Black has
advantage of
knight and
two pawns



white to move, but black has
advantage ($2 \times 1 + 3 = 5$)



white to move, can capture
queen, and take the advantage

For chess, typically **linear** weighted sum of **features**

$$Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

e.g., $w_1 = 9$ with

$f_1(s) = (\text{number of white queens}) - (\text{number of black queens}), \text{ etc.}$

pawn = 1
knight or bishop = 3
rook = 5
queen = 9
king = ∞

Minimax for heuristic evaluation functions

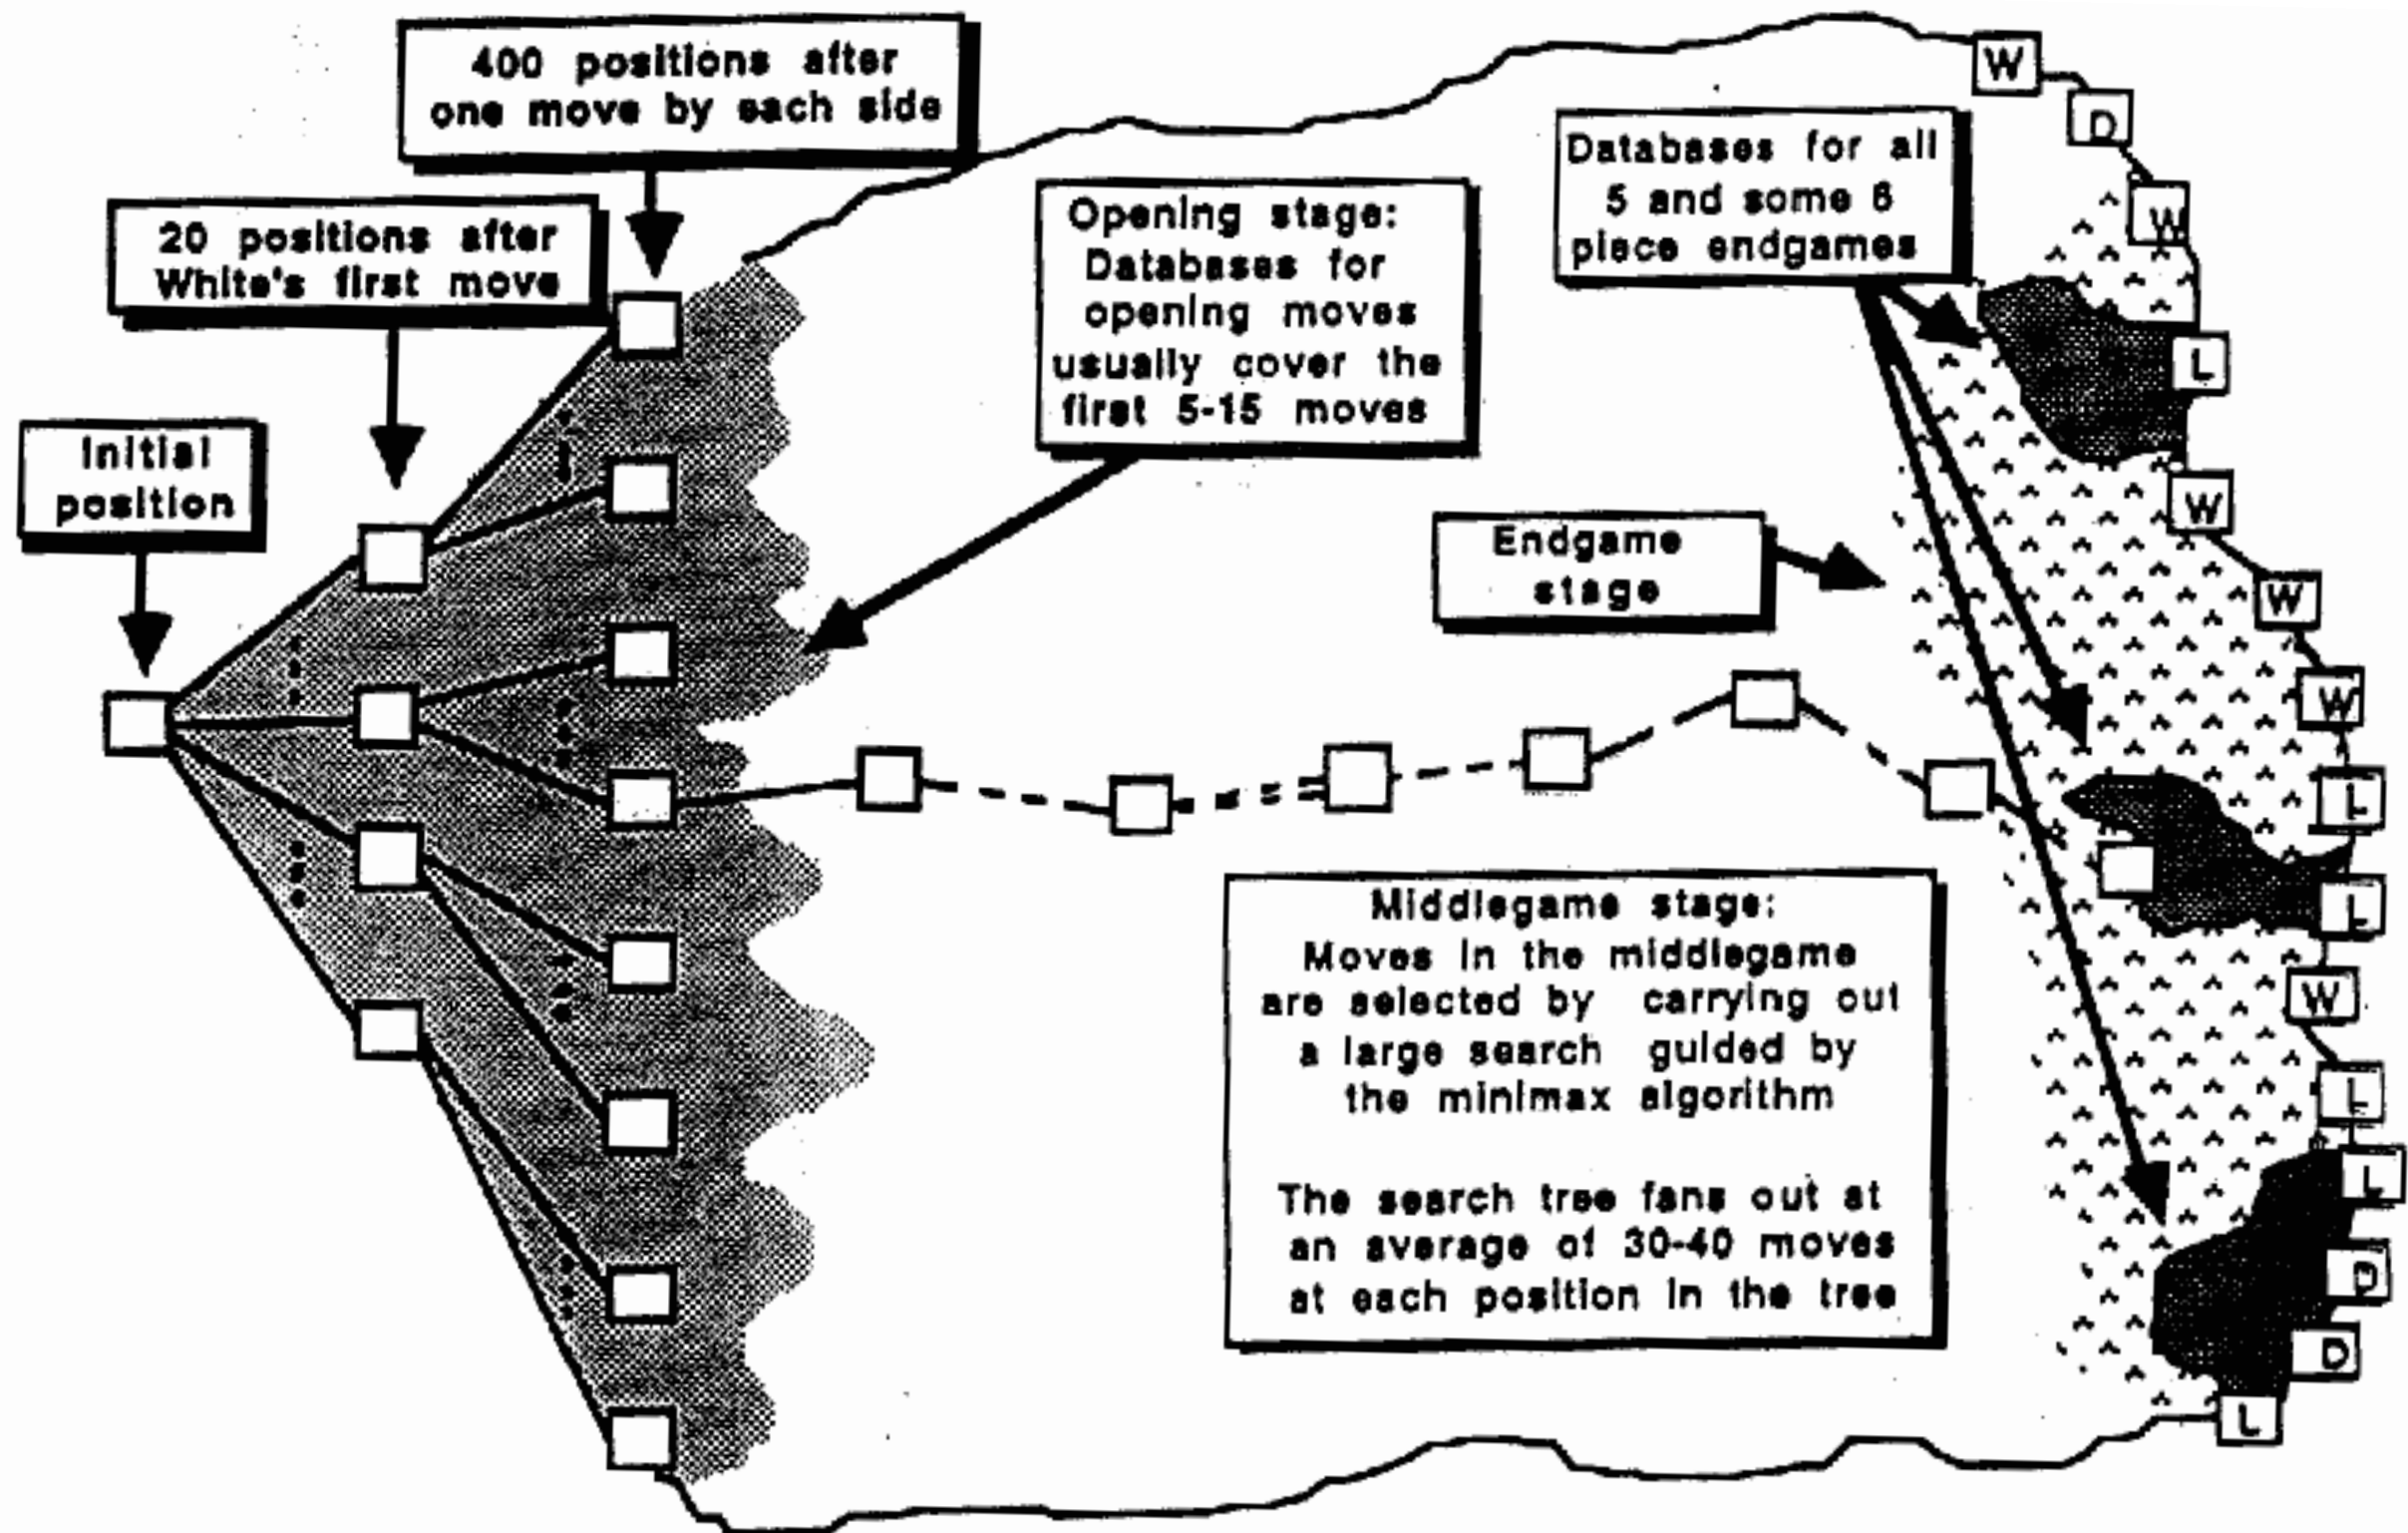
- Can't search very deep in complex games, need to cutoff search early
- Instead of using Utility of Terminal State, use heuristics at maximum depth.
- $H\text{-Minimax}(s,d) =$
 - $\text{Eval}(s)$ if $\text{cutoff-test}(s,d)$
 - $\max a \text{ in Actions}(s) \quad H\text{-Minimax}(\text{Result}(s,a), d+1)$ if $\text{Player}(s) = \text{Max}$
 - $\min a \text{ in Actions}(s) \quad H\text{-Minimax}(\text{Result}(s,a), d+1)$ if $\text{Player}(s) = \text{Min}$

A brief history of computer game playing

- computer considers moves (Babbage, 1846)
- algorithm for perfect game play (Zermelo, 1912; Von Neumann, 1944)
- finite horizon, approximate evaluation (Zuse, 1945; Wiener, 1948)
- first chess program (Turning, 1951)
- machine learning to improve evaluation accuracy (Samuel, 52-57)
- pruning to allow deeper search (McCarthy, 1956)

| | | |
|--|---|------|
| Shannon, Turing | minimax search with scoring function | 1950 |
| Kotok/McCarthy program, & ITEP program | alpha-beta | 1966 |
| Mac Kack | transposition tables | 1967 |
| Chess 3.0 - 4.9 | iterative-deepening depth-first search | 1975 |
| Belle | special purpose circuitry | 1978 |
| Cray Blitz | parallel search | 1983 |
| Hitech | parallel evaluation | 1985 |
| Deep Blue | parallel search and special-purpose circuitry | 1987 |
| many others since | | |

Computer Chess




Aspects of sophisticated evaluation functions

- Difference between player and opponent of:
 - # pieces
 - mobility
 - king position
 - bishop pair
 - rook pair
 - control of center (piecewise)
 - etc

Deep Blue

- Deep Blue evaluation function
 - ~6,000 different features
 - weighting depends on board configuration (downloaded into computer after each move)
- Weights are based on:
 - database of ~900 grandmaster games
 - tries to adjust weights to match play of grandmasters
- Also adjusted manually by Grand Master Joel Benjamin.



Player to move

value of knight's position in Deep Blue

Deep Blue's search

- ~200 million moves / sec
- 3.6×10^{10} moves in 3 minutes
 - 7 plies of uniform depth minimax search
 - 10-14 plies of uniform depth alpha-beta search
- software searches first
- specialized hardware searches last 5 ply

Deep Blue's hardware

- 32-node RS6000 SP multicomputer
- Each node
 - 1 IBM Power2 Super Chip
 - 16 chess chips
 - move generation
 - evaluation
 - some end-game heuristics and small end-game databases
- 32 GByte opening & endgame database

Role of computing power




Figure 6.23. Relationship between the level of play by chess programs and the size of the tree searched during a three minute move.

Deep Blue
defeats
Kasparov,
1997

Deep Fritz
defeats
Kramnik,
2006*

*but Kramnik
overlooked a
mate-in-one


Kasparov and Deep Blue



<http://www.youtube.com/watch?v=NJarxpYyoFI>

Kasparov *Business Insider* interview (May 24, 2017)

- Kasparov had just published *Deep Thinking*
 - reflections on AI over the 20 years after Deep Blue
- *Is AI a result or a process?*
 - *result*: Yes. Deep Blue is obviously intelligent because it plays chess at a grand master level
 - *process*: No. It's not human-like intelligence; just a fast "alarm clock."
- On making chess moves:
 - 1% calculation
 - 99% understanding of the game: patterns, previous experience
 - For a machine it's the exact opposite.



Mastering the game of Go with deep neural networks and tree search

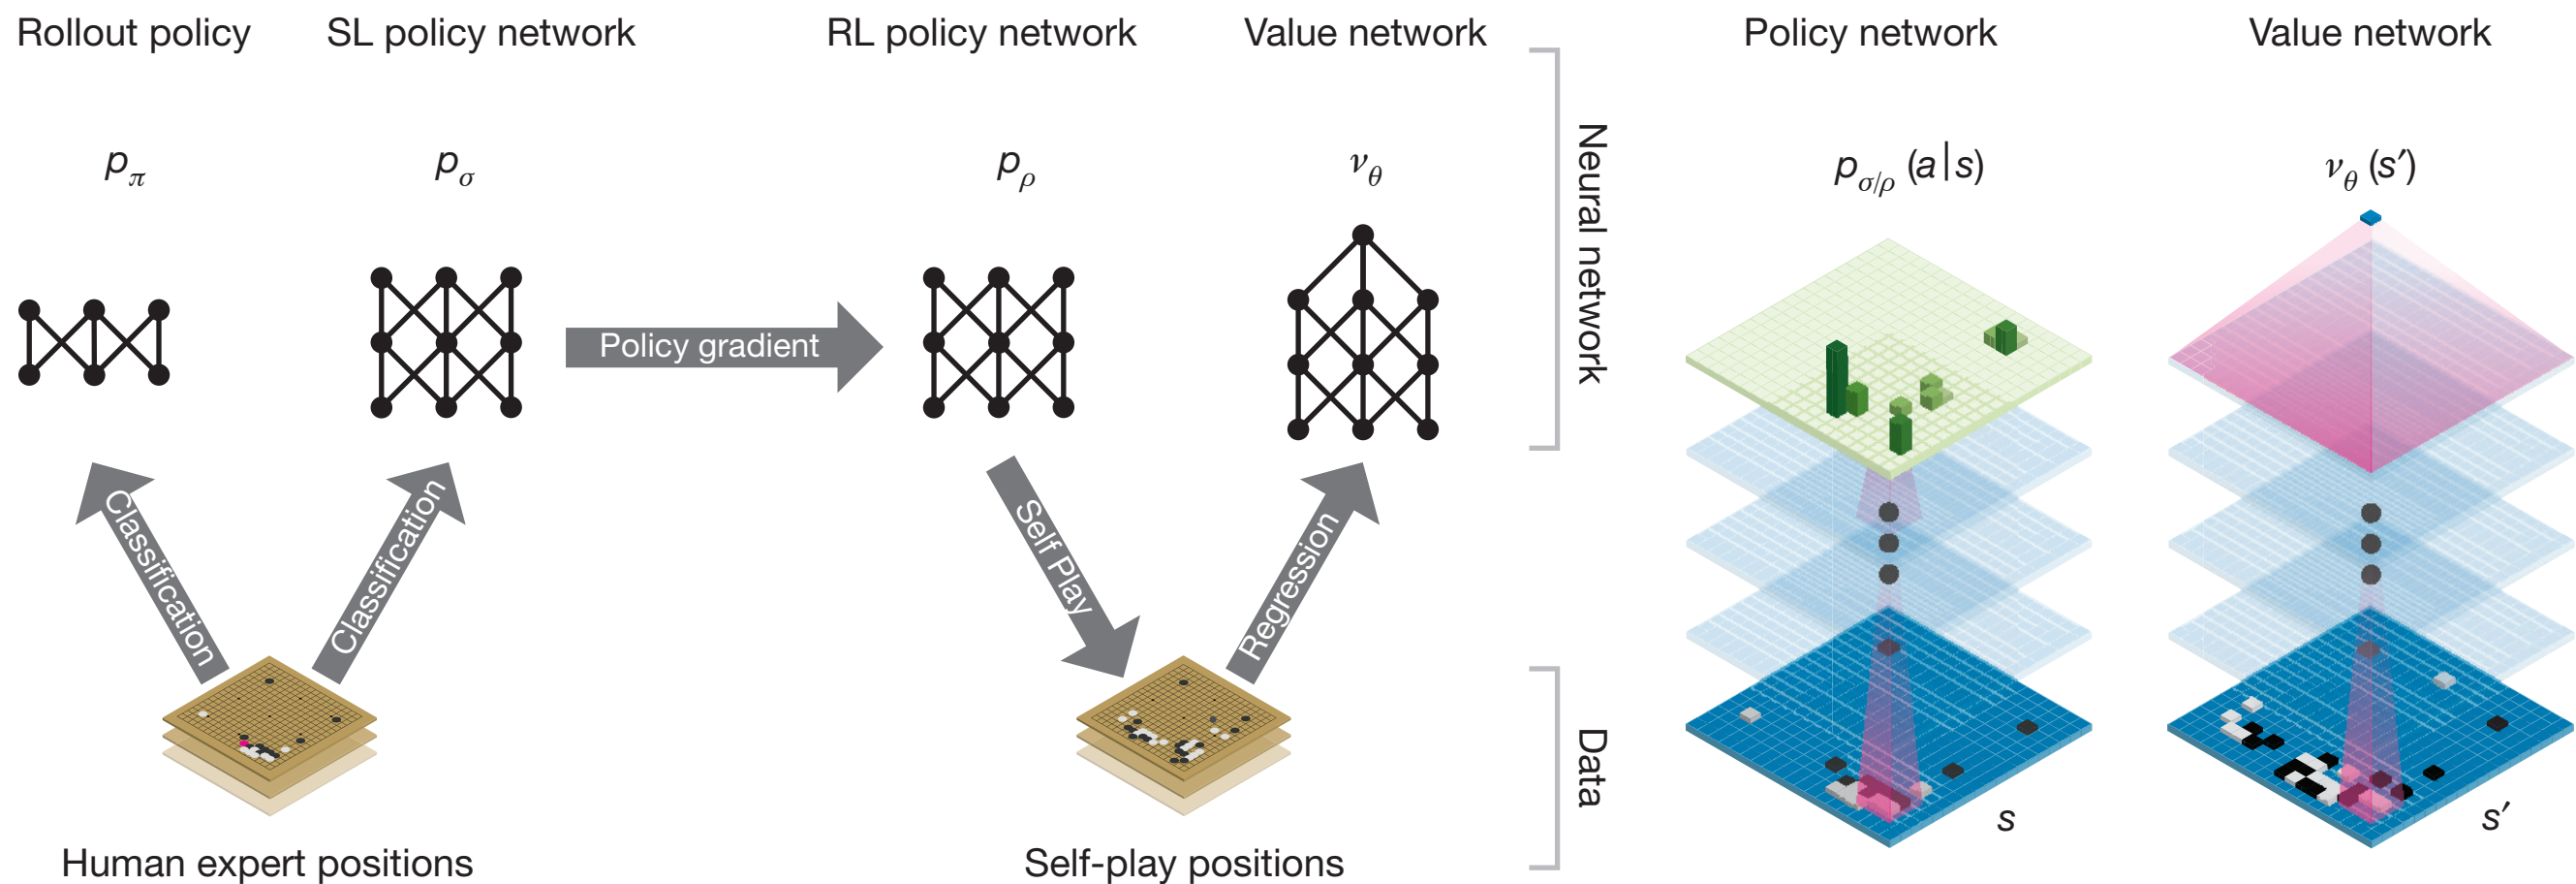
David Silver^{1*}, Aja Huang^{1*}, Chris J. Maddison¹, Arthur Guez¹, Laurent Sifre¹, George van den Driessche¹, Julian Schrittwieser¹, Ioannis Antonoglou¹, Veda Panneershelvam¹, Marc Lanctot¹, Sander Dieleman¹, Dominik Grewe¹, John Nham², Nal Kalchbrenner¹, Ilya Sutskever², Timothy Lillicrap¹, Madeleine Leach¹, Koray Kavukcuoglu¹, Thore Graepel¹ & Demis Hassabis¹

- Alpha Go: *Nature* 28 Jan 2016
- Combined many machine learning techniques (some of which we will cover in future lectures)
- By the end of the course, you should be able to understand the basic approach
- Beat 18-time world champion Lee Sedol in March 2016
- Won 4/5 games.



Alpha Go


- Key ideas:
 - *learn* what to do instead of using classical search-based techniques
 - use *value* (neural) networks to evaluate board positions
 - use *policy* networks to select moves
- No look-ahead search, but
 - played at the level of state-of-the-art Monte Carlo tree-search algorithms
 - achieved 99.8% winning rate against other Go programs



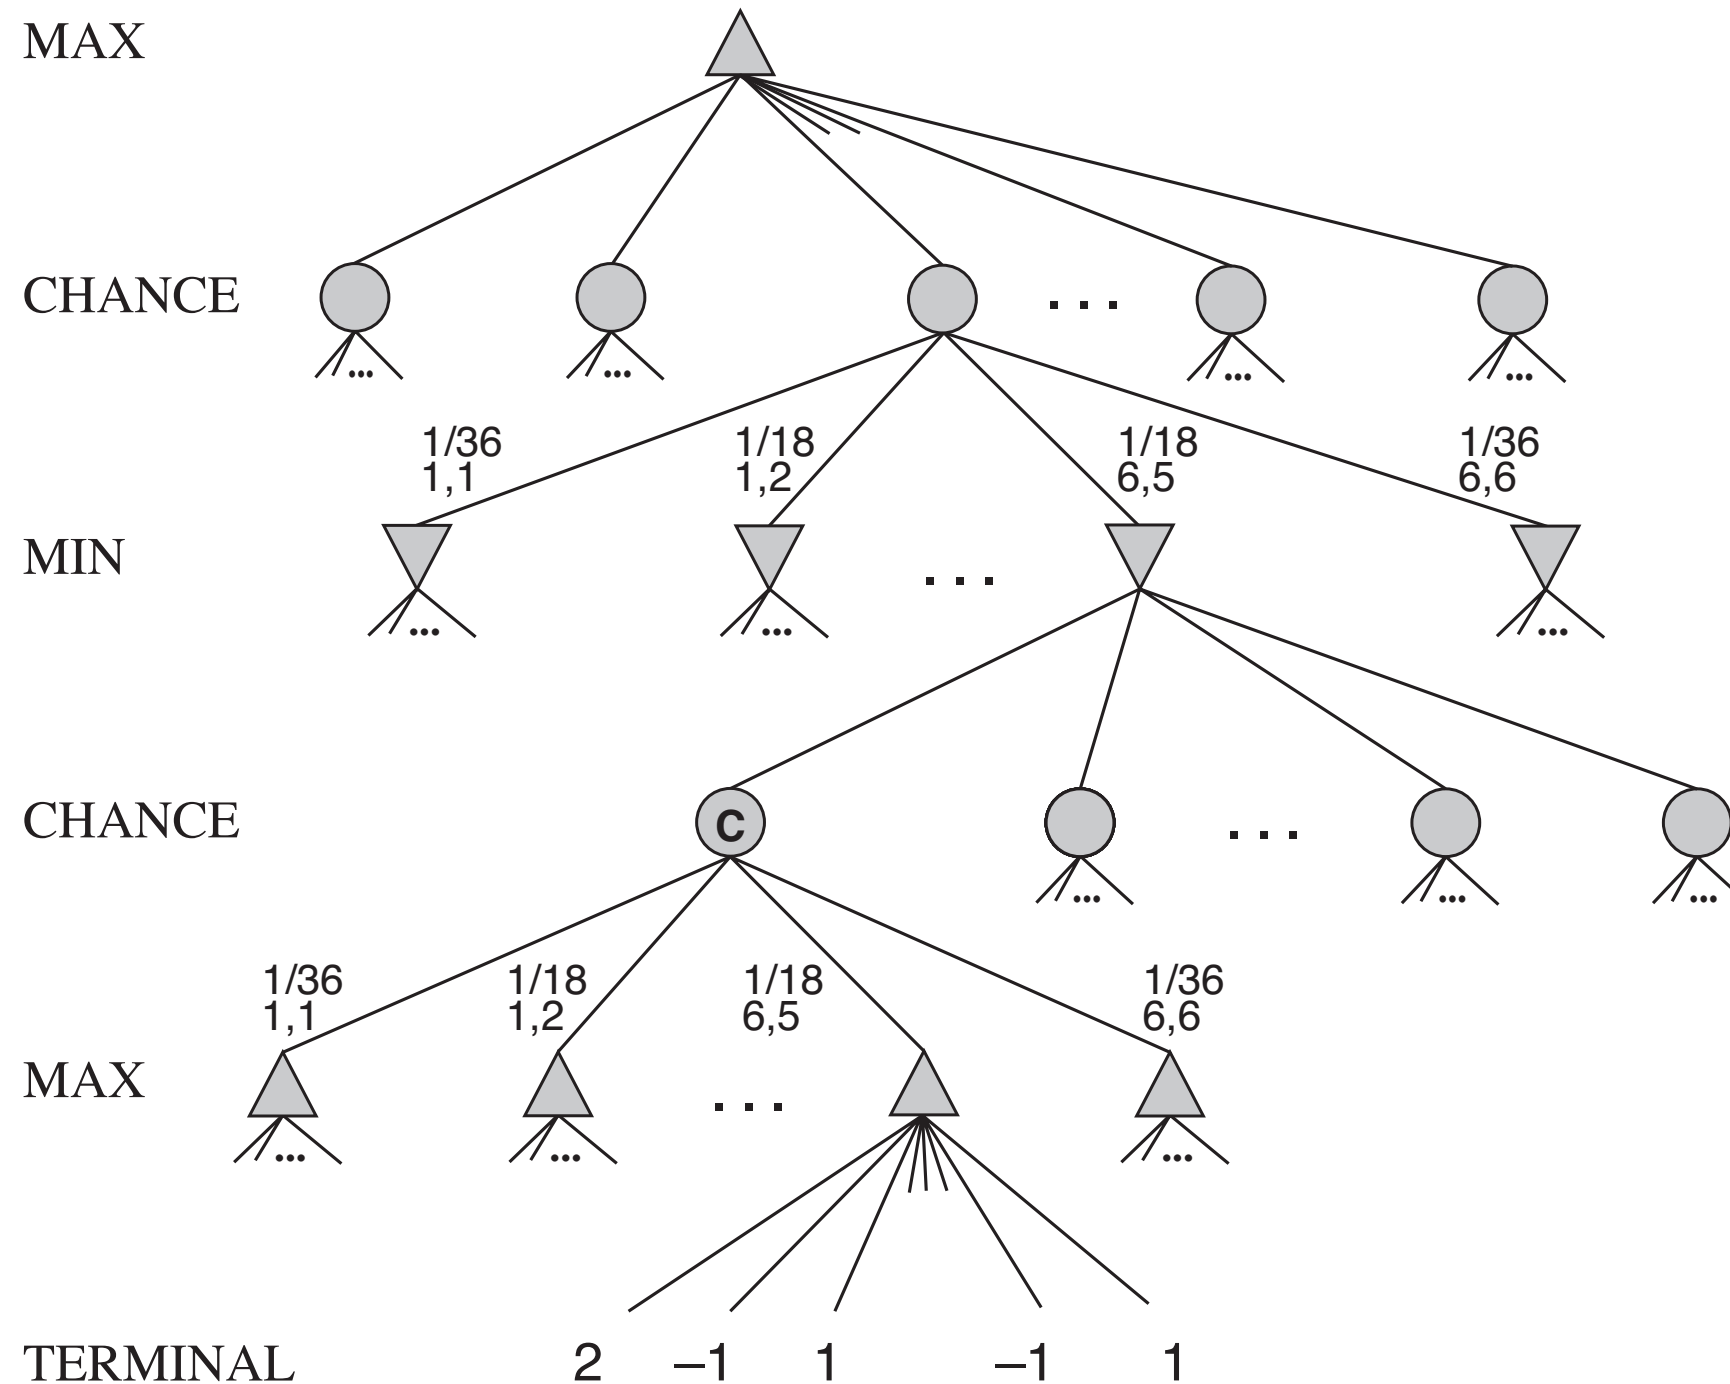
- Use supervised learning to predict moves of Human experts
- The policy network generates $p(a|s)$:
 - probability distribution over legal moves given current board state
- Use reinforcement learning to improve SL policy network
 - generate huge numbers of games for learning using self-play
- Train a value network to predict game outcome (probability of winning)

Stochastic games

- In backgammon, object is to move all pieces off the board.
- White moves clockwise, black counter-clockwise
- Die rolls determine possible moves
- A piece can move to any position, unless occupied by multiple opponents.
- If there is only one, that piece is captured and sent back to start.
- For a roll of 6 & 5, White has 4 legal moves:
(5-10,5-11)
(5-11,19-24)
(5-10,10-16),
(5-11,11-16)

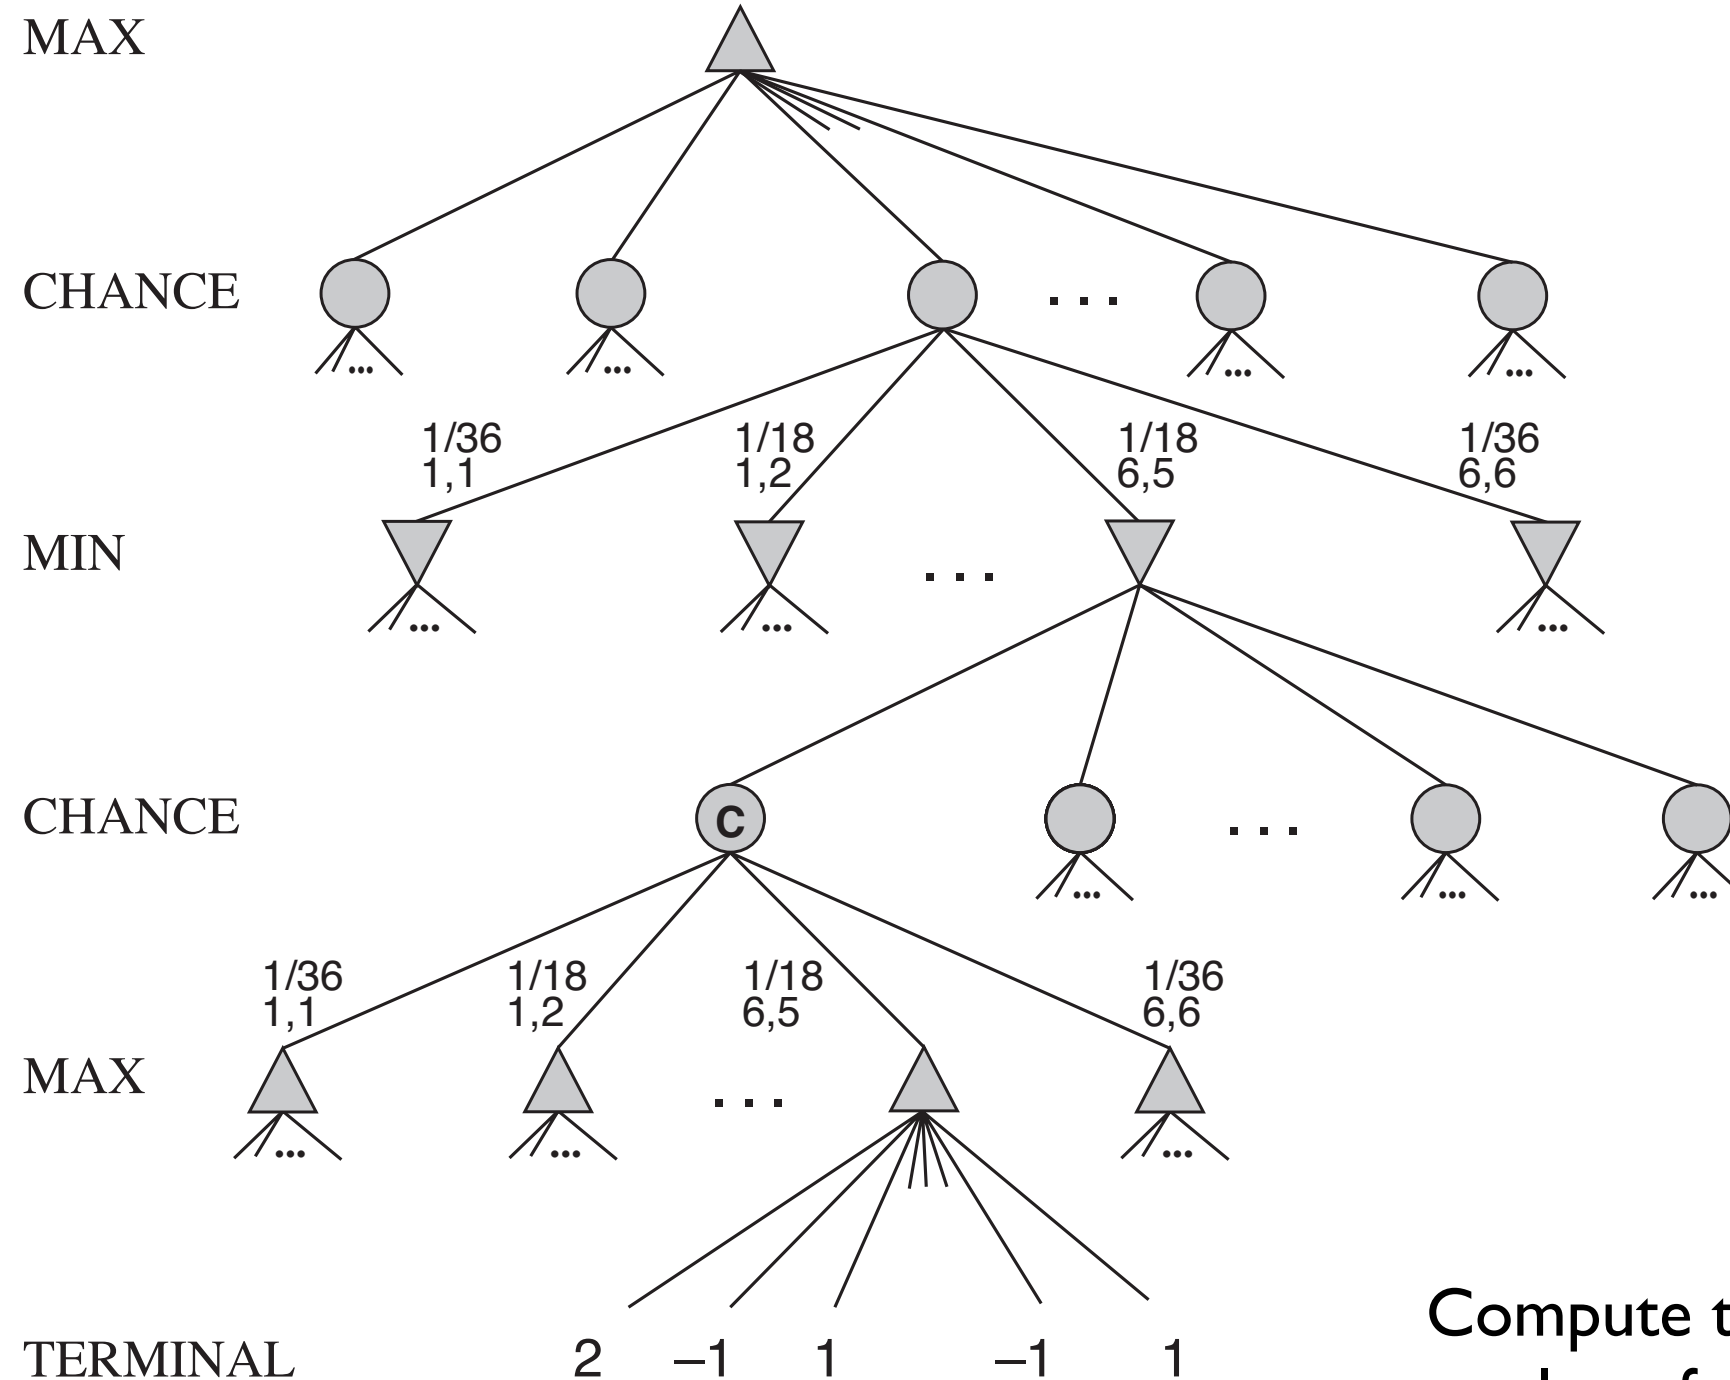


Stochastic game tree for backgammon position



How should we modify the evaluation function to accommodate chance?

Stochastic game tree for backgammon position



$$\text{eval}(s) = \sum_{s' \in \text{succ}(s)} p(s') \text{minimax}(s')$$

Properties of stochastic games

- die rolls: 21 possible rolls with two dies \Rightarrow *much* larger branching factor
- backgammon
 - ≈ 20 legal moves
 - considering opponent responses effective branching factor ≈ 400
- traditional searches are completely ineffective
- α - β pruning isn't very helpful
- Evaluation function becomes much more important
- TD-gammon (1992):
 - depth 2 search
 - learned a very good evaluation function:
a neural network, learned with temporal different learning
 - knowledge free: achieved intermediate-level play
 - when combined with expert-designed features: world champion level
 - stopped improving after 1.5M games